

Preparation

- **Log in to mirage**

```
$ ssh -l login mirage[0-2].ucar.edu
```

Use CryptoCard or Yubikey

- **Copy the example source files**

```
$ cp -r /glade/home/dnagle/Fortran-III .
```



Modern Fortran III for Computational Scientists

Consulting Services Group
Dan Nagle (Presenter)
May 24, 2012



Outline

- Documents
- Coarrays
- Trial Problems
- Implementations

Documents

- WG5 paper N1824.pdf
(www.nag.co.uk/sc22wg5)
- Rice University (caf.rice.edu)
- FDIS J3/10-007r1.pdf (www.j3-fortran.org)
- Modern Fortran Explained by
Metcalf, Reid, Cohen

PGAS

- Partitioned Global Address Space
- Examples include:
 - UPC (upc.gwu.edu)
 - Titanium (titanium.cs.berkeley.edu)
- Original Paper at <ftp://ftp.numerical.rl.ac.uk/pub/reports/nrRAL98060.ps.gz>



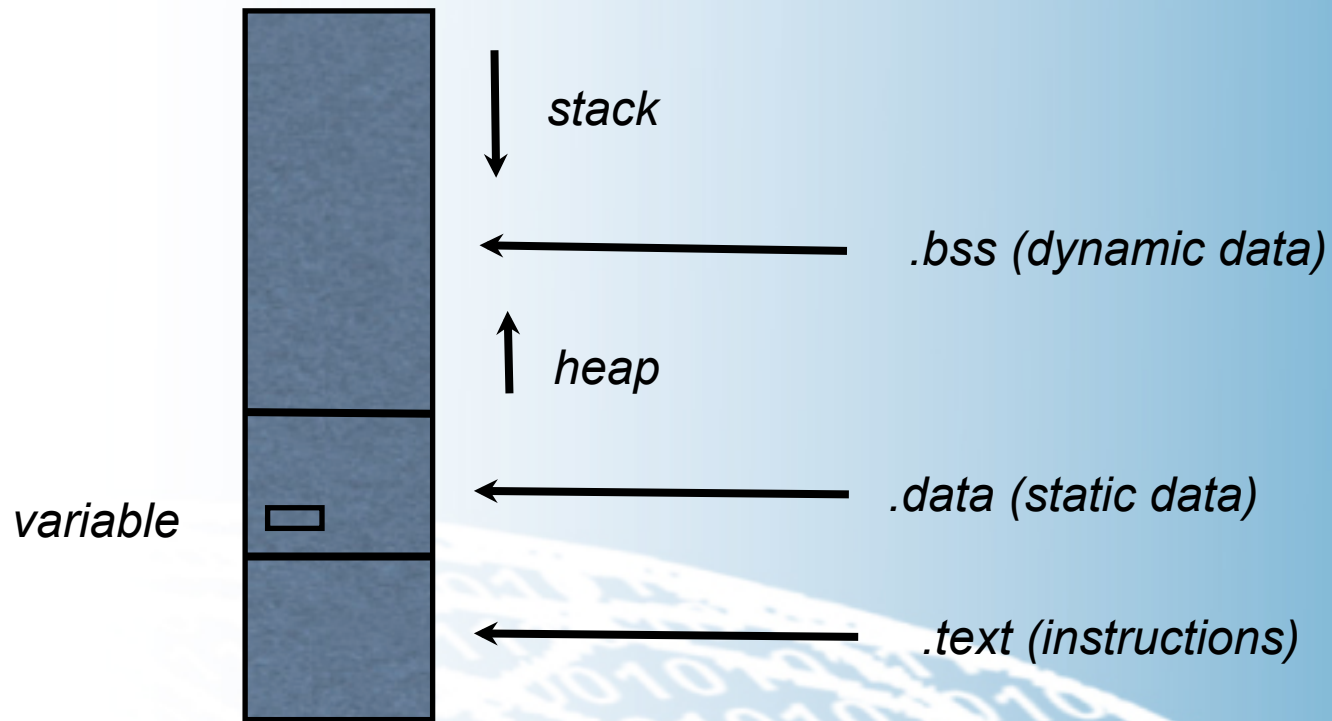
Brief History

- Agreed at Delft 2005
- Discussed at Fairfax 2006
- Again at London 2007
- Las Vegas Compromise 2008 (core & more)
- Tokyo Further Compromise 2008
- Finally Agreed Las Vegas 2010

Coarray Concepts

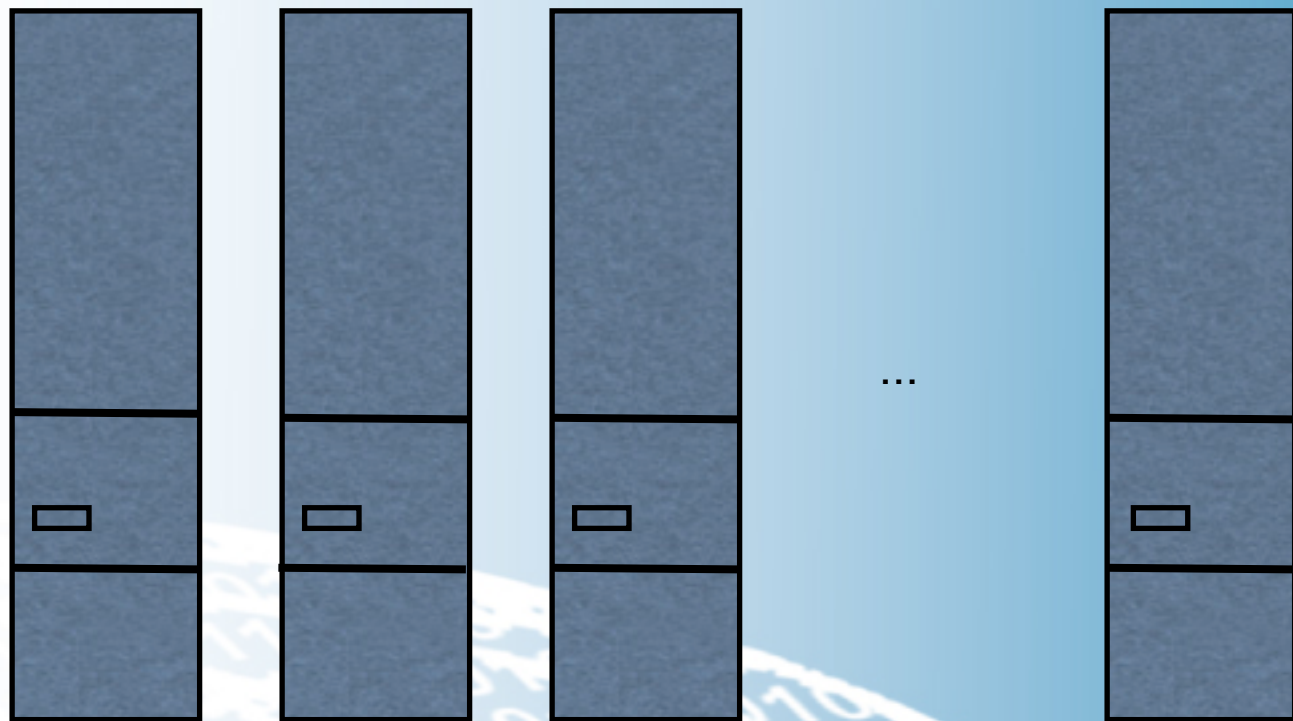
- A program is treated as if it were replicated at the start of execution, each replication is called an image.
- Each image executes asynchronously.
- A coarray is indicated by trailing [].
- A data object without trailing [] is local.
- Explicit synchronization statements are used to maintain program correctness.

Memory Basics (Single Image)



variable

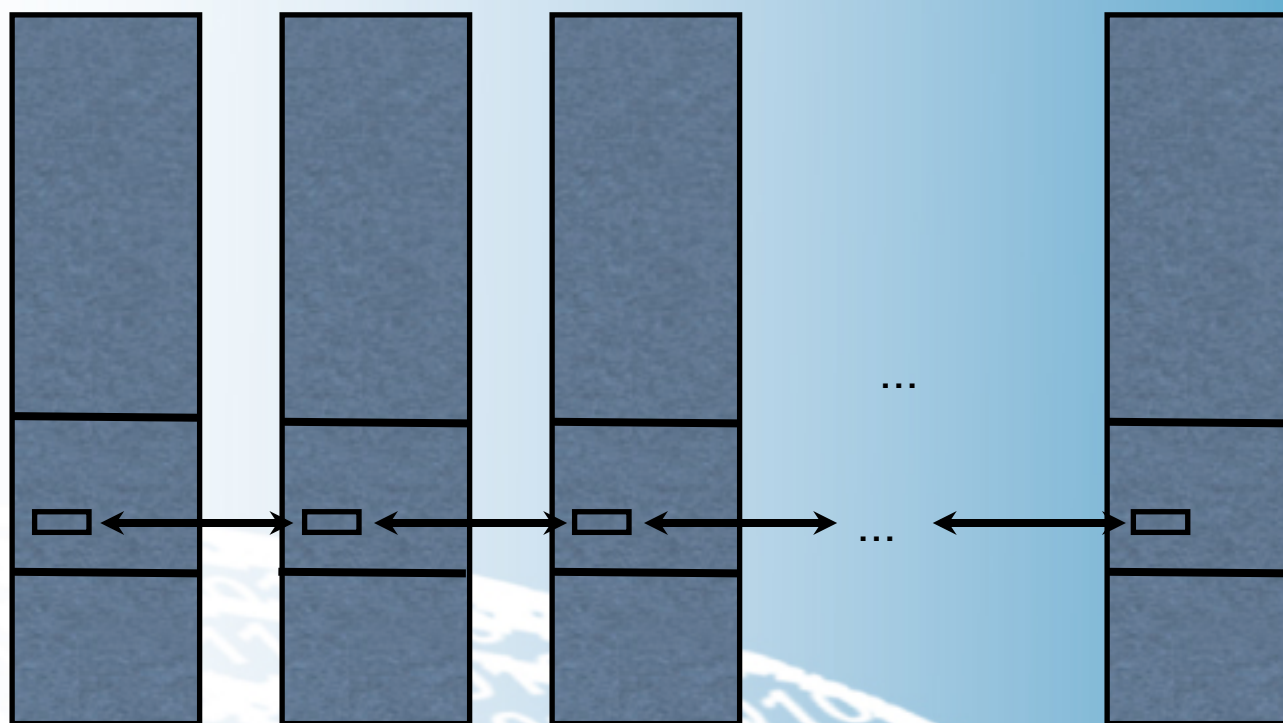
Memory Basics (Multiple Images)



Memory Basics

Coarrays

coarray



Memory Organization

- A static variable is stored at a fixed memory location.
- The fixed location is independent of the number of replications of the program.
- The coarray attribute connects the storage locations on each replication.
- The remote image may use the same address as the local image to access a data item.



Using Dynamic Coarrays

- A coarray may have the same address on all images.
- A coarray program must synchronize when a coarray is allocated or deallocated, or when an automatic local variable in a referenced subprogram is a coarray.

Image Indexes

- An image has an image index, that is a number between one and the number of images, inclusive.
- The images may be addressed as a rectangular array.
- The mapping from coindexes to image index is the same as an ordinary array.

Coarray Declarations

```
! a scalar coarray
```

```
real :: a[ *]
```

```
! an array coarray
```

```
real, dimension( n) :: ar[ *]
```

```
! a scalar coarray
```

```
real :: ca[ 10, 10, *]
```

```
! an array coarray
```

```
real :: caa( m, n)[ 10, 10, *]
```

Coarray Declarations

```
! a coarray with corank > 1
```

```
real :: a[ 10, *]
```

```
! sum of rank + corank <= 15
```

```
real, dimension( n) :: ar[ *]
```

```
! another scalar coarray
```

```
real :: ca[ -10: 10, *]
```

```
! the last coextent is always *
```


Simple Coarray Usage

! get a value from another image

a = b[4]

! send a value to another image

b[n] = a

! get an array value

c(1: n) = ca(1: n)[i]

! put an array value

caa(1: n)[j] = aa(1: n)



Coarrays as Actual Arguments

- ! a coarray may be used
- ! as an ordinary actual argument
- ! may have intent in, out, in out

```
call my_sub( b[ 4 ])
```

```
x = f( y[ j ])
```



Coarrays as Dummy Arguments

- ! a coarray dummy argument
- ! needs an explicit interface
- ! intent may be in, out, in out

```
subroutine my_sub( b)  
real, intent( in out) :: b[ *]
```

```
pure function func( y)  
real, intent( in) :: y( :) [ *]
```



Segments

- A segment is a piece of code between synchronization points.
- A compiler is free to apply all its optimizations within a segment.
- Segments are ordered by synchronization statements or dynamic memory actions.
- Segments may be ordered or unordered.

Synchronization

! synchronize all the images

sync all

! synchronize with a set of images

sync images(this_image() + 1)

sync images(list_images(i: j))

*sync images(*)*

! wait for memory quiet

sync memory

Coarray Local Variables

```

! function has coarray local variables
function f( x)
..
    real :: local_a[ *]
..
end function f

! reference f
! this causes two synchronizations:
! one at the call, one at the return
y = f( x1)
    
```



Allocatable Coarrays

! declaring allocatable coarrays

```
real, allocatable, dimension( :) :: ac[ :]
```

! allocate as usual

! this causes a synchronization

```
allocate( ac( n)[ *], stat= ... )
```

! deallocate as usual

! this causes a synchronization

```
deallocate( ac, stat= ... )
```



Derived Types with Coarrays

! a derived type with a coarray component

```
type :: co_array_t
    real, dimension( :) :: dtc[ *]
end type co_array_t
```

! a variable of type co_array_t

! this must be a scalar

```
type( co_array_t) :: my_dt_coarray
```

! a coarray of derived type

! my_type has no coarray components

```
type( my_type), dimension( :) :: dtc[ *]
```



Inheriting Derived Types with Coarrays

! a derived type with a coarray component

```
type :: parent_t
    real, dimension( n) :: dtc[ *]
end type parent_t
```

! inherit from parent_t

! if child_t has coarray components,

! so must the parent_t

```
type, extends( parent_t) :: child_t
    real, dimension( n) :: adtc[ *]
end type child_t
```



Derived Types with Coarrays

! a derived type with a coarray component

```
type :: has_coarray_t
    real, dimension( :) :: dtc[ *]
end type has_coarray_t
```

! a variable of type co_array_t

! this must be a scalar

```
type( has_coarray_t) :: my_dt_coarray
```

! the coarrayness exists at only one level

Derived Types with Coarrays

! derived type without coarray components

```
type :: no_coarray_t
    real :: vx, vy, vz
end type no_coarray_t
```

! a coarray of derived type

! no_coarray_t has no coarray components

```
type( no_coarray_t) :: dtc( n)[ *]
```

! the coarrayness exists at only one level

Coarrays & Pointers

```

! a coarray cannot be a pointer
! so make a derived type
type :: cptr
  real, dimension( :), pointer :: p
end type cptr
type( cptr) :: a, b[ *]
! all components assigned
a = b
! a% p is undefined
a = b[ j]

```



Coarray Intrinsic

```

! number of images
nim = num_images()
! my image number
me = this_image()
! cobounds
l = lcobound( ca, dim= 1)
u = ucobound( ca, dim= 2)
! index of cosubscripts
i = image_index( ca, subs)

```



Locks

```
! declare a lock coarray
type( lock_type) :: work_lock[ *]

! lock a critical resource
lock( work_lock)
...
unlock( work_lock)

! lock a neighbor's lock
lock( work_lock[ me + 1])
...
unlock( work_lock[ me + 1])
```



Critical Section

! one image at a time

critical

! update the critical resource

i_crit = i_crit + 1

! end one image at a time

end critical

Atomic Intrinsics

```
integer( atomic_int_kind) :: iflag[ *]
logical( atomic_logical_kind) :: al[ *]
```

```
!   define iflag on image i to be 42
call atomic_define( iflag[ i], 42)
```

```
!   get value of al on image j
call atomic_ref( local_flag, al[ j])
```

Simple Example

```
program small_hello_world
  integer :: me
  continue

  me = this_image()
  write( unit= *, fmt= '( a, i0, a)' ) &
    'Image ', me, ' says "Hello, world!"

  stop 'normal completion'
end program small_hello_world
```



Monte Carlo Toy

- Using four images, compute an approximation of π as four times the ratio of points within a circle to points within the containing square.
- Image one should read the number of trials per image and print the sum of all the image's estimates.
- Remember to seed the random number generator differently on each image.



$M \times M$

- Image one reads two $N \times N$ matrices.
- Decide upon a storage scheme among the four images, and distribute the matrix.
- Decide upon a computation scheme (that is, a blocking scheme), and multiply the two matrices.
- Image one writes the product matrix.



Digits Home

- Image one reads 16 four-digit numbers, where each digit is in the range [1-4].
- Distribute one fourth of the numbers to each of four images.
- Shuffle the numbers to where the ones digit is equal to the image number where it is stored.
- Repeat for the tens, hundreds, thousands digits.
- Print from each image at each step.

3d-fft

- Image one reads a rank-three array.
- Distribute to each of four images.
- Compute a three dimensional FFT in place.
- Image one prints the result.





Reduction

- Each image computes a value from its image number (for example, the image number squared).
- Sum the computed numbers in a logarithmic number (of images) of steps.
- Without a broadcast, the sum is on each image.
- Verify the correctness of the sum on each image, and print the sum from image one.



Implementations

- Cray
- IBM (PPC-only, Beta-only)
- g95 (but is it maintained?)
- Intel (12.1+)
- gfortran (4.6 single-image only)

Future Coarrays

- John Reid's Summary N1824
- Bill Long's Draft of the TS N1858 = 11-176
 - Pre Garching
- Metcalf, Reid, & Cohen, Modern Fortran Explained
 - http://www.amazon.com/Explained-Numerical-Mathematics-Scientific-Computation/dp/0199601429/ref=sr_1_1?s=books&ie=UTF8&qid=1326999081&sr=1-1

Basic Ideas

- Collective functions
 - Changed between 08-131r1 and 11-256r2
- Teams
 - Original team proposal versus Rice U proposal
- Notify/Query – How Much Like Events ?
- Parallel I/O – Features and Limitations ?
- Global Data Structures – coscalars or pointers ?
- Atomic Operations

Documents with Suggestions

- 08-131r1 (the original list of deferred features)
- 10-166 (Bill Long's 2010 draft of the Further Coarrays TS based on 08-131r1)
- N1835 (John Reid's 2010 summary list)
- N1856 (Rice U CAF group 2011 list)
- N1883 Comments on list (post Garching)
- 11-256r2 (John Reid's 2011 summary list) **



NCAR





The Size of the Coarray Additions

- Competing desires:
 - to manage the workload on compiler suppliers
 - to provide useful tools to applications programmers as quickly as possible
- What's missing ?
 - Must judge without extensive application programmer experience with f08 coarrays
- “Useful” is application-dependent

Collective Procedures

- Which ones ?
 - Similar to the existing reduction intrinsics ?
 - Similar to MPI reduction procedures ?
- Synchronous or Asynchronous ?
 - If synchronous, what affect on performance ?
 - If asynchronous, how ?
 - Signal completion via event variables ?

The Original Set

- From 08-131r1 (see also 07-007r3)
 - co_all, co_any, co_count
 - co_maxloc, co_maxval
 - co_minloc, co_minval
 - co_product, co_sum
 - co_findloc ?
- Most Had Arg Lists (source, result [, team])
- Synchronous (Image Control Statements)

The Current Favorites

- As per 11-256r2
- co_bcast, co_max, co_min, co_reduce, co_sum
- asynchronous ?
- copy_async ?
- max-and-copy, min-and-copy ?
- apply to non-coarrays ?

Teams

- Teams were originally lists of image indexes
- Teams might be subdivided from a parent team via the scheme proposed by the Rice U CAF group
- What should teams do ? (synchronize and ...)
 - label blocks ? (the **with team** construct)
 - procedures ? (intrinsic or external ?)
 - I/O ? (**team=** on control list)

Parallel Input/Output

- Sequential Access or Direct Access or Both ?
 - If sequential, control record order ?
 - If direct, control record access ?
 - how to manage races ?
- Linked with Teams how ?
 - **team=** on control list ?

Notify/Query

- As is (identified by image number), or as first class events (identified by event variables) ?
- Synchronous or queued ?
- Applied to a team or global ?
- What about libraries ?
 - Do library writers prefer event variables ?

If First Class Events, Details ?

- Apply to collective intrinsics ?
 - copy_async: Rice U CAF proposal has 3 events: source ready, destination ready, copy ready
- Apply to asynchronous I/O ?
- Apply to any other wait operation ?
- Do events queue or signal ?
 - If they queue, how to distinguish individual events ?

Global Data Structures

- Coscalars (Germany)
- Copointers (Rice U CAF)
 - and cotarget
- Pointer attribute allowed on coarrays (IBM)

Any other business

- Asynchronous copy (Rice U CAF)
- More atomic operations (Cray)
 - compare-and-swap is likely
 - the rest of Cray's current set ?
 - add, and, or, xor
 - and corresponding fetch-and-op versions
- Support for irregular grids ?
 - affinity for teams ?

The Process Needs Input

- In February, set the initial feature list to start the discussion for the International meeting
 - might be viewed as the opening US position
- The June International meeting at IBM Toronto Labs will set the work list
- What are your concerns ?
- Discussion at <http://sea.ucar.edu/forums/updates-fortran-2008>



Modern Fortran III for Computational Scientists

Thanks for Attending!

May 24, 2012